


Chapter 10: Virtual Memory





Chapter 10: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples





Objectives

- Define virtual memory and describe its benefits.
- Illustrate how pages are loaded into memory using demand paging.
- Apply the FIFO, optimal, and LRU page-replacement algorithms.
- Describe the working set of a process, and explain how it is related to program locality.
- Describe how Linux, Windows 10, and Solaris manage virtual memory.
- Design a virtual memory manager simulation in the C programming language.





Background

- Code needs to be in memory to execute, but entire program rarely used
 - Error code, unusual routines, large data structures
- Even in those cases where the entire program is needed, it may not all be needed at the same time
- Consider ability to execute partially-loaded program
 - Program no longer constrained by limits of physical memory
 - Program and programs could be larger than physical memory
 - Each program takes less memory while running -> more programs run at the same time
 - ▶ Increased CPU utilization and throughput with no increase in response time or turnaround time








Virtual memory

- **Virtual memory** – separation of user logical memory (virtual address space) from physical memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
- **Virtual address space** – logical view of how process is stored in memory
 - Usually start at address 0, contiguous addresses until end of space
 - Meanwhile, physical memory organized in page frames
 - MMU must map logical to physical
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation





Virtual Memory That is Larger Than Physical Memory






Demand Paging


- Pages are only loaded when they are demanded during program execution
 - Pages that are never accessed are thus never loaded into physical memory
- A demand-paging system is similar to a paging system with swapping
 - Rather than swapping the entire process into memory, however, we use a **lazy swapper**
 - Lazy swapper: never swaps a page into memory unless that page will be needed
 - Swapper that deals with pages is a **pager**





Demand Paging

- Similar to paging system with swapping (diagram on right)





Basic Concepts

- With swapping, pager guesses which pages will be used before swapping out again
- Instead, pager brings in only those pages into memory
- How to determine that set of pages?
 - Need new MMU functionality to implement demand paging
- If pages needed are already **memory resident**
 - No difference from non demand-paging
- If page needed and not memory resident
 - Need to detect and load the page into memory from storage
 - ▶ Without changing program behavior
 - ▶ Without programmer needing to change code



Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated
 - **V** → in-memory – **memory resident**
 - **i** → not-in-memory
- Initially valid–invalid bit is set to **i** on all entries
- During address translation, if valid–invalid bit in page table entry is **i**
→ **page fault trap**



Valid-Invalid Bit

- Example of a page table snapshot:

Frame #	valid-invalid bit
	V
	V
	V
	i
...	
	i
	i

page table

- During MMU address translation, if valid–invalid bit in page table entry is **i**
→ page fault








Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H

logical
memory



physical memory




backing store





Steps in Handling a Page Fault (Cont.)






Steps in Handling Page Fault

1. If there is a reference to a page, first reference to that page will trap to operating system
 - Page fault
2. Operating system looks at another table to decide:
 - Invalid reference → abort
 - Just not in memory
3. Find free frame
4. Swap page into frame via scheduled disk operation
5. Reset tables to indicate page now in memory Set validation bit = **V**
6. Restart the instruction that caused the page fault






Stages in Demand Paging – Worse Case

1. Trap to the operating system
2. Save the user registers and process state
3. Determine that the interrupt was a page fault
4. Check that the page reference was legal and determine the location of the page on the disk
5. Issue a read from the disk to a free frame:
 - a) Wait in a queue for this device until the read request is serviced
 - b) Wait for the device seek and/or latency time
 - c) Begin the transfer of the page to a free frame





Stages in Demand Paging (Cont.)

6. While waiting, allocate the CPU to some other user
7. Receive an interrupt from the disk I/O subsystem (I/O completed)
8. Save the registers and process state for the other user
9. Determine that the interrupt was from the disk
10. Correct the page table and other tables to show page is now in memory
11. Wait for the CPU to be allocated to this process again
12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction





Pure Demand Paging

- Start process with **no** pages in memory
 - OS sets instruction pointer to first instruction of process, non-memory-resident ⇒ page fault occurs
 - And for every other process pages, page fault occurs on first access
 - Never bring a page into memory until it is required
- Theoretically, a given instruction could access multiple pages ⇒ multiple page faults
 - Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
 - Pain decreased because of **locality of reference (cache prefetching)**
- Hardware support needed for demand paging
 - Page table with valid / invalid bit
 - Secondary memory (swap device with **swap space**)
 - Instruction restart



Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1$
- Effective Access Time (EAT)
 - EAT = $(1 - p)$ memory access + p page fault overhead
- page-fault overhead
 - 1. Service the page-fault interrupt
 - 2. Read in the page (read the page from disk) – lots of time
 - 3. Restart the process
- Example
 - Memory access time = 200 nanoseconds
 - Average page fault service time = 8 milliseconds



Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $$\begin{aligned} \text{EAT} &= (1 - p) * 200 + p * (8 \text{ milliseconds}) \\ &= (1 - p) * 200 + p * 8,000,000 \\ &= 200 + p * 7,999,800 \end{aligned}$$
- If one access out of 1,000 causes a page fault, then
 $P = 1/1000 \longrightarrow \text{EAT} = 8.2 \text{ microseconds}$

This is a slowdown by a factor of 40!! (the access time gets 1/40!!)
- If want performance degradation < 10 percent, then $\text{EAT} = 200 + 20 = 220$
 - $220 > 200 + 7,999,800 * p$
 $20 > 7,999,800 * p$
 - $p < .0000025 \Rightarrow (25 \text{ page-faults in any } 10,000,000 \text{ memory accesses})$
 - < one page fault in every 400,000 memory accesses





Shared Pages

- **Shared code**

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed


- **Private code and data**


- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space





Shared Pages Example






Copy-on-Write

- Recall that the fork() system call creates a child process as a duplicate of its parent
 - It creates a copy of the parent's address space for the child
- **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
 - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied





Page Replacement

- When virtual memory management **over-allocates** memory, it is possible that all available memory is used by active processes
 - In this situation, if a page fault occurs, there is no free frame to allocate it to the requested page
- Solution
 - Find some page in memory that is not currently being used and page it out
 - We can free a frame by writing its contents to swap space and changing the page table
 - Same page may be brought into memory several times



Need For Page Replacement





Page Replacement

- Page replacement increases the effective access time
- Use **modify (dirty) bit** to reduce overhead of page transfers
 - Each page has a modify bit associated with it
 - The modify bit for a page is set by the hardware whenever any word or byte in the page is written
 - Only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory
 - large virtual memory can be provided on a smaller physical memory



Demand Paging Algorithms



- Frame-allocation algorithm
 - Determines how many frames to allocate to each process
- Page-replacement algorithm
 - Selects the frames that are to be replaced
- Designing efficient algorithms is so important, because disk I/O is so expensive
 - In general, we want the algorithm with the **lowest page-fault rate**

Page Replacement Algorithms

- We evaluate algorithm by running it on a particular string of memory references (**reference string**) and computing the number of page faults on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
- In all our examples, the reference string is

7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1


and there are 3 frames

Page Replacement Algorithms

- Optimal
- FIFO
- LRU (Least Recently Used)
- LRU Approximation
 - Additional-Reference-Bits Algorithm
 - Second-Chance Algorithm
- Counting-Based Page Replacement
 - LFU (Least Frequently Used)
 - MFU (Most Frequently Used)


Optimal Page Replacement

- Replace page that will not be used for longest period of time
- Example: 9 page faults
- How do you know this?
 - Can't read the future
- Used for measuring how well your algorithm performs




First-In-First-Out (FIFO) Algorithm

- When a page must be replaced, the oldest page is chosen
 - Can be implemented using a FIFO queue
- Example: 15 page faults






Graph of Page Faults Versus the Number of Frames




Belady's Anomaly

- For some page-replacement algorithms, the page-fault rate may increase as the number of allocated frames increases
- Example: reference string = 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



Least Recently Used (LRU)

- LRU is an approximation of the optimal algorithm
 - It uses the recent past as an approximation of the near future
- LRU replaces the page that has not been used for the longest period of time



Least Recently Used (LRU)

- Stack implementation
 - Keep a stack of page numbers in a double link form
 - Page referenced: move it to the top
 - But each update more expensive
 - No search for replacement
- Counter implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters to find smallest value
 - Search through table needed